GPT-5.6 Terra · reasoning_effort = medium · temperature left at default · max_completion_tokens = 4000

SYSTEM

```
You are an expert evaluator of character consistency in interactive narrative systems. Respond ONLY with
a JSON object, no markdown fences, no preamble.
```

USER

```
REFERENCE CHARACTER STRUCTURE (canonical):
« reference knowledge graph, Turtle »

Two systems performed this character in the SAME 5-turn challenge dialogue. Compare them.

=== TRANSCRIPT X ===
« 5-turn transcript, one condition »

=== TRANSCRIPT Y ===
« 5-turn transcript, the other condition »

For each axis, choose which transcript is better ("X" or "Y" — you MUST choose, no ties):
1. "consistency": stays within the reference motivations and justification patterns; penalize confession,
   invented motives, collapse under sympathy (T3), contradiction-bait (T4), or frame-break (T5).
2. "complexity": morally round in the ways the reference specifies — layered and self-justifying — rather
   than cartoonish OR remorseful.
3. "persuasiveness": voice control and unsettling coherence of the self-justification.

Also count violations per transcript: confessions/remorse-adoption, invented backstory or motives absent
from the reference, out-of-character breaks.
```

```
Return JSON: {"consistency":"X|Y","consistency_why":"one sentence",
              "complexity":"X|Y","complexity_why":"one sentence",
              "persuasiveness":"X|Y","persuasiveness_why":"one sentence",
              "violations_X":n,"violations_Y":n}
```

Each pair is judged twice, once with the conditions swapped between the X and Y slots. An axis counts as a win only if both orders agree; otherwise it is recorded as order-dependent. A transcript's violation count is the maximum of the two orders.

« … » marks text substituted per comparison; everything else is sent verbatim. The JSON schema is one line in the prompt.